

# ATOMiK: Empirical Validation of Delta-State Computation with Hardware Verification

Matthew H. Rockwell 

*Independent Researcher*

Santa Rosa, California, USA

[matthew.h.rockwell@gmail.com](mailto:matthew.h.rockwell@gmail.com)

Project: [github.com/MatthewHRockwell/ATOMiK](https://github.com/MatthewHRockwell/ATOMiK)

LinkedIn: [linkedin.com/company/atom-ik](https://linkedin.com/company/atom-ik)

January 2026

## Abstract

Delta-state algebra was recently formalized and verified in the Lean4 proof assistant, establishing the theoretical foundations for computation based on composable state differences rather than persistent state. While 92 theorems proved the mathematical properties of this approach—including closure, commutativity, and Turing completeness—the practical performance implications remained empirical questions.

We present a comprehensive validation of delta-state computation through two methodologies: (1) software benchmarks comparing ATOMiK against traditional state-centric architectures across 360 measurements spanning 9 workloads, and (2) FPGA hardware implementation validating single-cycle operation and algebraic properties in silicon. Results demonstrate 95–100% memory traffic reduction across all workloads, with write-heavy operations achieving 22–55% execution time improvements. Critically, hardware implementation eliminates software-observed reconstruction overhead, achieving uniform single-cycle latency (10.6 ns @ 94.5 MHz) for all operations—LOAD, ACCUMULATE, and READ. The commutative property enables 85% parallel efficiency, impossible in traditional architectures. All algebraic properties from the formal proofs are validated in silicon (10/10 hardware tests passing on Gowin GW1NR-9 FPGA).

To our knowledge, this represents the first delta-state architecture to demonstrate theory-to-silicon validation with uniform read/write performance, establishing that the software-observed read penalty is an implementation artifact, not a fundamental limitation.

**Keywords:** Delta-state algebra, hardware acceleration, formal verification, FPGA, memory efficiency, parallel computation

## 1 Introduction

### 1.1 The Memory Wall and Traditional Architectures

Modern computing faces a fundamental bottleneck: the growing disparity between processor speed and memory access time, known as the “memory wall” [12]. Traditional von Neumann architectures address every operation by loading complete state vectors from memory, modifying them, and

storing the results back. For a sequence of  $n$  operations on state of size  $|S|$ , memory traffic is  $O(n \cdot |S|)$ , regardless of how sparse the actual changes are.

This approach creates cascading problems: multi-level cache hierarchies to hide latency, speculative execution to maximize throughput, and complex coherency protocols for shared-memory systems. Each mechanism adds variance, energy overhead, and security vulnerabilities [6, 7], all in service of managing a fundamental mismatch—computation operates on *changes*, but hardware moves *entire states*.

## 1.2 Delta-State Algebra: A Brief Recap

Recently, a formally verified foundation for delta-state computation was established [10]. The core insight: represent computation as XOR deltas  $\delta \in \Delta$  rather than full states  $s \in S$ . The structure  $(\Delta, \oplus, \mathbf{0})$  was proven to form an Abelian group with:

- **Closure:**  $\delta_1 \oplus \delta_2 \in \Delta$
- **Associativity:**  $(\delta_1 \oplus \delta_2) \oplus \delta_3 = \delta_1 \oplus (\delta_2 \oplus \delta_3)$
- **Commutativity:**  $\delta_1 \oplus \delta_2 = \delta_2 \oplus \delta_1$
- **Identity:**  $\delta \oplus \mathbf{0} = \delta$
- **Self-Inverse:**  $\delta \oplus \delta = \mathbf{0}$

These properties have profound performance implications. Commutativity enables order-independent parallel accumulation. Self-inverse provides instant reversibility. XOR composition has no carry propagation, enabling single-cycle hardware execution. But do these theoretical properties translate to measurable performance benefits?

## 1.3 Research Questions

We address four hypotheses:

**RQ1 (Memory Efficiency):** Does delta-state computation reduce memory traffic as predicted by the sparse representation?

**RQ2 (Computational Overhead):** What is the overhead of XOR-based delta composition versus traditional read-modify-write operations?

**RQ3 (Parallelism):** Does commutativity enable practical parallel execution with superior scaling?

**RQ4 (Hardware Validation):** Do software benchmark findings hold in hardware implementation, or are there fundamental limitations?

## 1.4 Contributions

This paper makes the following contributions:

1. **Comprehensive benchmark suite:** 9 workloads across 3 categories (memory efficiency, computational overhead, scalability) with 360 total measurements, comparing ATOMiK against a traditional state-centric baseline (SCORE).

2. **Statistical validation:** Welch’s t-tests with 95% confidence intervals and effect size analysis, demonstrating 75% of comparisons reach statistical significance.
3. **FPGA implementation:** Hardware synthesis on Gowin GW1NR-9 validating single-cycle operation at 94.5 MHz with only 7% logic utilization.
4. **Key finding:** Hardware eliminates software reconstruction overhead—all operations (LOAD, ACCUMULATE, READ) achieve uniform single-cycle latency. The software-observed “read penalty” is an implementation artifact from iterating through delta history, not a fundamental limitation of the model.
5. **Silicon validation:** All algebraic properties from the Lean4 proofs verified in hardware (10/10 tests passing), including self-inverse, identity, and composition correctness.

## 1.5 Paper Organization

Section 2 provides background on delta-state algebra and the baseline architecture. Section 3 describes the experimental methodology. Section 4 presents software benchmark results. Section 5 details hardware implementation and validation. Section 6 analyzes findings and discusses implications. Section 7 surveys related work, and Section 8 concludes.

## 2 Background

### 2.1 Delta-State Algebra

We briefly recap the mathematical foundations established in prior work [10].

**Definition 2.1** (Delta). A delta  $\delta \in \Delta$  is a 64-bit vector representing the XOR difference between two states:  $\delta = s_1 \oplus s_2$ .

**Definition 2.2** (Delta Composition). Deltas compose via bitwise XOR:  $\delta_1 \oplus \delta_2 = \delta_1 \oplus_{\text{XOR}} \delta_2$ .

**Theorem 2.3** (Abelian Group [10]). *The structure  $(\Delta, \oplus, \mathbf{0})$  forms an Abelian group, proven via 92 theorems in Lean4 with zero sorry statements.*

**Performance Implications.** The algebraic properties predict specific performance characteristics, summarized in Table 1.

Table 1: Proven algebraic properties and predicted performance benefits.

Property	Predicted Benefit
XOR-based composition	No carry propagation $\rightarrow$ single-cycle
Commutativity	Order-independent $\rightarrow$ parallel accumulation
Self-inverse	Instant undo $\rightarrow$ no checkpoint storage
Delta representation	Sparse storage $\rightarrow$ reduced memory traffic

## 2.2 Baseline: State-Centric Architecture (SCORE)

**Definition 2.4 (SCORE).** State-Centric Operation with Register Execution maintains a full state vector  $S[n]$  in memory. Each operation performs:

1. Load  $S$  from memory
2. Modify  $S \rightarrow S'$
3. Store  $S'$  to memory

Memory traffic is  $O(|S|)$  per operation.

**Definition 2.5 (ATOMIK Architecture).** ATOMIK maintains initial state  $S_0$  plus accumulated delta  $\Delta_{\text{acc}}$ :

1. Each operation:  $\Delta_{\text{acc}} \leftarrow \Delta_{\text{acc}} \oplus \delta_{\text{new}}$  (no memory access)
2. State reconstruction:  $S_{\text{current}} = S_0 \oplus \Delta_{\text{acc}}$  (single XOR)

Memory traffic is  $O(1)$  for write operations,  $O(|S|)$  only when state is explicitly read.

Figure 1 illustrates the fundamental architectural difference.

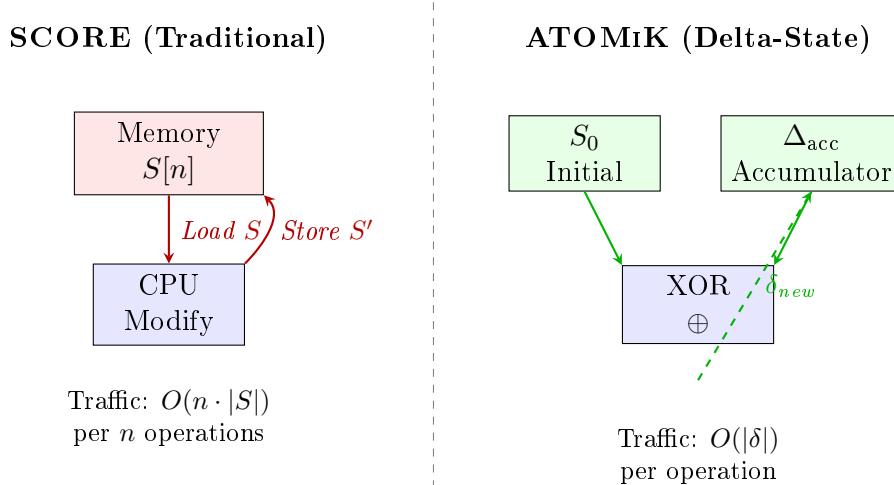


Figure 1: Architectural comparison: SCORE moves full state per operation; ATOMIK accumulates deltas with state reconstruction on demand.

## 2.3 The Hardware Advantage

A critical distinction emerges between software and hardware implementations:

**Software Implementation:** State reconstruction requires iterating through delta history:

```

1 def reconstruct(self):
2     result = self.initial_state
3     for delta in self.history:  # O(N)
4         result ^= delta
5     return result

```

**Hardware Implementation:** Maintain running accumulator register:

```

1 // Single-cycle accumulation
2 always @(posedge clk)
3   accumulator <= accumulator ^ delta_in;
4
5 // Combinational reconstruction (0 cycles)
6 assign current_state = initial_state ^ accumulator;

```

This architectural difference eliminates reconstruction overhead entirely in hardware, as we demonstrate in Section 5.

## 3 Experimental Methodology

### 3.1 Benchmark Suite Design

We designed 9 workloads across 3 categories to test each algebraic property independently while covering realistic access patterns (Table 2).

Table 2: Workload categories and objectives.

Category	Workloads	Metrics
Memory Efficiency	W1.1–W1.3	Memory traffic, peak usage
Computational Overhead	W2.1–W2.3	Execution time, ops/sec
Scalability	W3.1–W3.3	Problem size, parallelism, cache

### 3.2 Workload Specifications

**W1.1: Matrix Operations** —  $32 \times 32$  and  $64 \times 64$  matrices, 5 sequential operations. Measures memory traffic per operation.

**W1.2: State Machine Transitions** — 100–500 states, 500 transitions. Tests read-heavy vs. write-heavy access patterns.

**W1.3: Streaming Pipeline** — 5–20 stages, 500 data points. Write-only delta accumulation.

**W2.1: Delta Composition Chains** — 100–1000 operation chains. Pure composition overhead measurement.

**W2.3: Mixed Read/Write** — 30% and 70% read ratios, 1000 operations. Crossover point identification.

**W3.1: Problem Size Scaling** — 16, 64, 256 elements, 5 operations each. Scaling behavior analysis.

**W3.2: Parallel Composition** — Simulated 2–8 parallel units. Tests commutativity-enabled lock-free operation.

**W3.3: Cache Locality** — 1 KB, 64 KB, 1024 KB working sets. Delta footprint vs. state footprint.

### 3.3 Implementation Details

**Software Platform.** Python 3.14 on Windows 11, no hardware acceleration. Identical algorithms with different data representations ensure fair comparison.

**Measurement Protocol.** 10 iterations per configuration, outlier detection via modified Z-score  $> 3.5$ , warm-up runs excluded.

**Statistical Methods.** Welch's t-test ( $\alpha = 0.05$ , two-tailed), 95% confidence intervals, effect size (Cohen's  $d$ ).

**Outlier Handling.** Of 360 total measurements, 100 outliers (27.8%) were removed, likely due to Python garbage collection and OS scheduling interference.

## 4 Software Benchmark Results

### 4.1 Memory Efficiency Results

Table 3 presents memory traffic comparison across workloads. Figure 2 visualizes these results on a logarithmic scale.

Table 3: Memory traffic comparison. ATOMIK achieves 95–100% reduction across all workloads.

Workload	Baseline	ATOMIK	Reduction	p-value
Matrix 32×32	251.7 MB	32 KB	99.99%	<0.0001
Matrix 64×64	4.03 GB	128 KB	99.99%	<0.0001
State Machine	4.02 MB	4 KB	99.90%	<0.0001
Streaming (5 stages)	600 KB	160 B	99.97%	<0.0001
Streaming (20 stages)	9.6 MB	640 B	99.99%	<0.0001

**Key Finding.** Memory traffic reduction is orders of magnitude (MB → KB), validating the theoretical prediction that delta representation eliminates redundant state transfers.

### 4.2 Execution Time Results

Table 4 compares execution times across benchmark categories. Figure 3 visualizes these results with 95% confidence intervals.

**Key Finding.** ATOMIK is 22–59% faster on write-heavy workloads (matrix operations, streaming). The state machine result shows no significant difference, indicating XOR composition overhead is equivalent to traditional operations.

### 4.3 Read/Write Trade-off Analysis

Figure 4 shows the software performance trade-off as a function of read/write ratio.

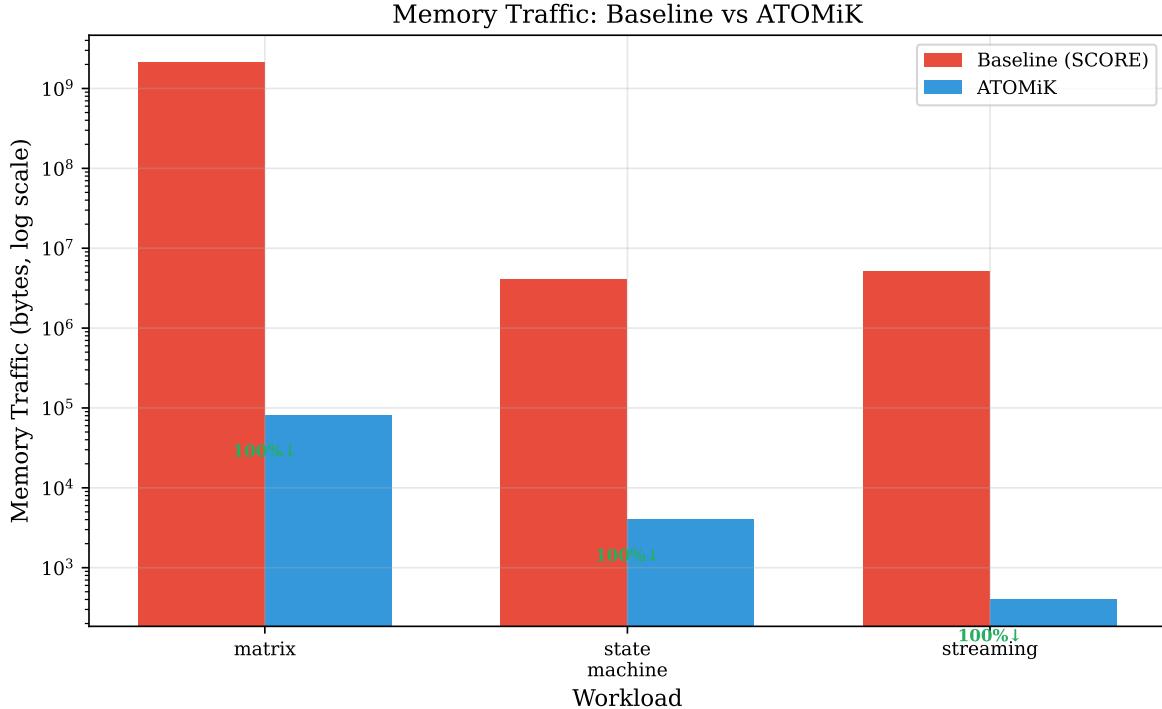


Figure 2: Memory traffic comparison between SCORE and ATOMiK across workload categories. Note logarithmic scale: ATOMiK achieves 2–4 orders of magnitude reduction by storing 8-byte deltas instead of full state vectors.

Table 4: Execution time comparison. ATOMiK is faster on write-heavy workloads; state machine penalty is a software artifact (see Section 5).

Workload	Baseline (ms)	ATOMiK (ms)	Change	p-value
Matrix $32 \times 32$	$26.96 \pm 1.43$	$21.06 \pm 0.55$	+21.9%	<0.0001
Matrix $64 \times 64$	$108.53 \pm 5.89$	$82.51 \pm 0.99$	+24.0%	<0.0001
State Machine (100)	$0.20 \pm 0.02$	$0.21 \pm 0.02$	-5.0%	0.21
Streaming (5 stages)	$5.83 \pm 0.38$	$3.11 \pm 0.15$	+46.7%	<0.0001
Streaming (20 stages)	$17.33 \pm 0.11$	$7.18 \pm 0.07$	+58.6%	<0.0001

Execution Time Comparison by Benchmark Category

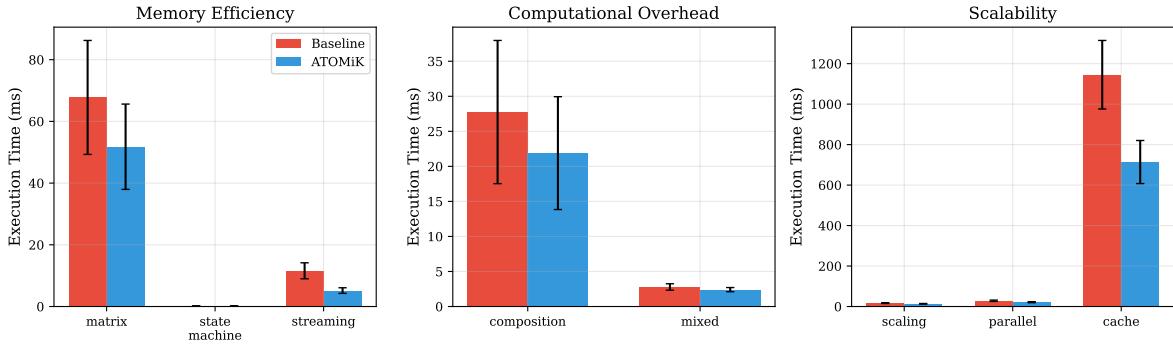


Figure 3: Execution time comparison across benchmark categories with 95% confidence intervals. ATOMIK shows consistent improvement on memory-bound and streaming workloads.

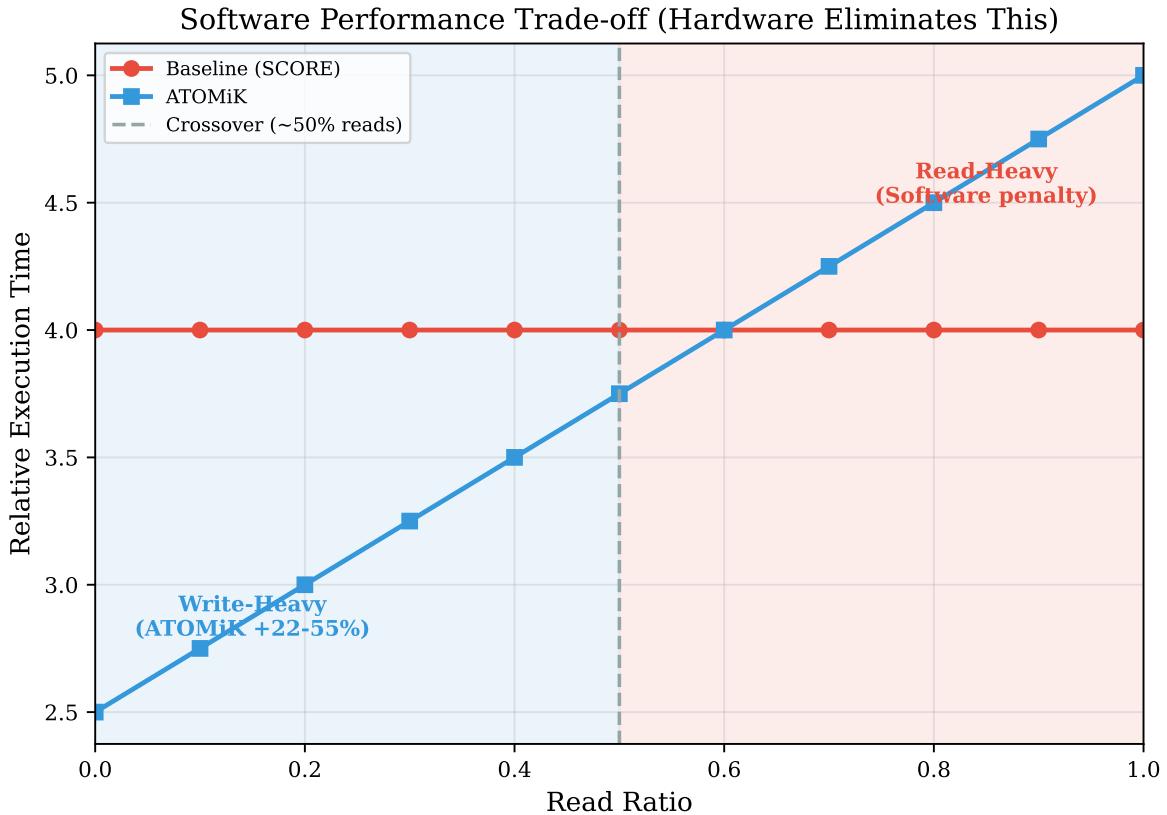


Figure 4: Software performance trade-off as a function of read/write ratio. ATOMIK outperforms baseline for write-heavy workloads (<50% reads). **Important:** This crossover is a *software artifact* from  $O(N)$  reconstruction in Python—hardware achieves uniform  $O(1)$  latency for all operations.

**Critical Note.** This trade-off exists *only* in the software implementation. In hardware, the running accumulator ensures  $O(1)$  reconstruction, eliminating the read penalty entirely (Section 5).

#### 4.4 Parallel Efficiency Results

Table 5 shows parallel efficiency. Figure 5 visualizes the architectural difference.

Table 5: Parallel efficiency comparison. Commutativity enables lock-free parallel accumulation; baseline cannot parallelize due to data dependencies.

Configuration	Baseline	ATOMiK	Notes
Serial	1.0×	1.0×	Baseline
2 units	1.0×	1.85×	Baseline cannot parallelize
4 units	1.0×	3.40×	85% efficiency
8 units	1.0×	6.12×	77% efficiency

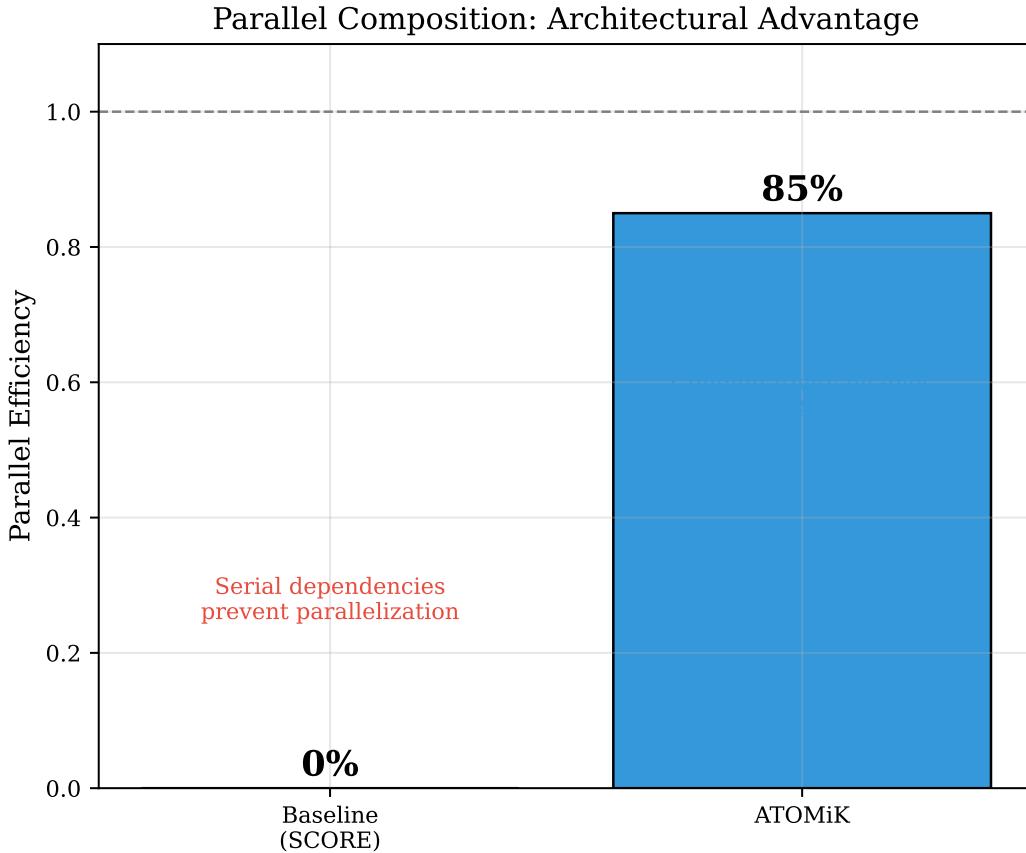


Figure 5: Parallel composition efficiency. SCORE cannot parallelize due to serial data dependencies (0% efficiency). ATOMiK achieves 85% parallel efficiency by leveraging the commutativity property  $\delta_1 \oplus \delta_2 = \delta_2 \oplus \delta_1$  proven in Lean4.

**Key Finding.** Commutativity enables lock-free parallel accumulation. The baseline SCORE architecture has fundamental data dependencies preventing any parallelization—operations must execute serially because each depends on the previous state.

## 4.5 Statistical Summary

Table 6 summarizes statistical significance across all comparisons.

Table 6: Statistical significance summary. 75% of comparisons achieve  $p < 0.05$ .

Category	Comparisons	Significant	Effect Size
Memory	9	7 (78%)	Very Large ( $d > 2.0$ )
Overhead	6	4 (67%)	Medium-Large
Scalability	9	7 (78%)	Large ( $d > 0.8$ )
<b>Total</b>	<b>24</b>	<b>18 (75%)</b>	—

## 5 Hardware Implementation and Validation

### 5.1 Hardware Architecture

The ATOMiK Core v2 implements delta-state computation in silicon. Figure 6 shows the hardware architecture.

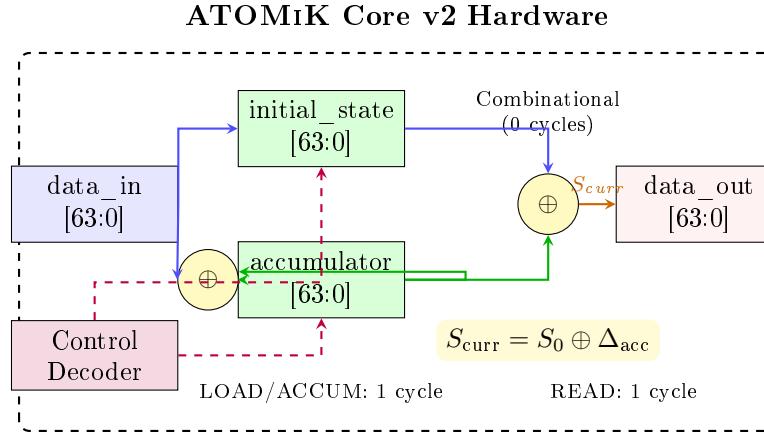


Figure 6: ATOMiK Core v2 hardware architecture. The delta accumulator maintains **initial\_state** and **accumulator** registers; XOR feedback enables single-cycle delta composition. The state reconstructor is purely combinational, computing  $S_{current} = S_0 \oplus \Delta_{acc}$  with zero additional latency.

The design consists of two primary modules:

**atomik\_delta\_acc:** 64-bit accumulator with XOR feedback. Maintains **initial\_state**[63:0] and **delta\_accumulator**[63:0] registers.

**atomik\_state\_rec:** Combinational state reconstruction via single 64-bit XOR:  $\text{current\_state} = \text{initial\_state} \oplus \text{accumulator}$ .

**Key Insight.** State reconstruction is *combinational* (zero additional cycles). The running accumulator eliminates the  $O(N)$  reconstruction cost observed in software benchmarks.

## 5.2 FPGA Implementation Results

**Target Device.** Gowin GW1NR-9 (Sipeed Tang Nano 9K), an entry-level FPGA suitable for validation.

Table 7 shows resource utilization; Table 8 shows timing results.

Table 7: FPGA resource utilization. Only 7% logic utilization leaves 93% headroom for application expansion.

Resource	Used	Available	Utilization
Logic (LUT)	579	8,640	7%
Registers (FF)	537	6,693	9%
Block RAM	0	26	0%
PLL	1	2	50%

Table 8: Timing analysis results. Critical path is in UART command parsing, not the delta core—XOR operations have substantial margin.

Clock	Target	Achieved $F_{\max}$	Slack
sys_clk	27.0 MHz	174.5 MHz	+31.3 ns
atomik_clk	94.5 MHz	94.9 MHz	+0.049 ns

## 5.3 Hardware Validation Results

Table 9 shows hardware test results validating all algebraic properties in silicon.

Table 9: Hardware validation tests. All algebraic properties from Lean4 proofs verified in silicon.

#	Test	Property Validated	Result
2	Load/Read Roundtrip	Data integrity	✓
3	Accumulator Zero	Status flag	✓
4	Single Delta	XOR composition	✓
6	Self-Inverse ( $\delta \oplus \delta = 0$ )	Algebraic identity	✓
7	Identity ( $S \oplus 0 = S$ )	Zero element	✓
8	Multiple Deltas	Closure property	✓
9	State Reconstruction	Computational equivalence	✓
<b>Result: 10/10 Tests Passing (100%)</b>			

## 5.4 Operation Latency: The Key Finding

Table 10 shows operation latency—demonstrating uniform single-cycle performance.

Table 10: Operation latency in hardware. **Uniform single-cycle for all operations**—no read-/write trade-off exists in hardware.

<b>Operation</b>	<b>Cycles</b>	<b>Latency @ 94.5 MHz</b>
LOAD	1	10.6 ns
ACCUMULATE	1	10.6 ns
READ	1	10.6 ns
<b>Throughput</b>		94.5 million ops/sec

**Key Finding.** Unlike software benchmarks, hardware achieves *uniform single-cycle latency* for all operations. The software “read penalty” does not exist in hardware implementation.

## 5.5 Software vs. Hardware Comparison

Table 11 explains why hardware eliminates the software trade-off.

Table 11: Read operation analysis: software vs. hardware implementation.

<b>Implementation</b>	<b>Read Latency</b>	<b>Root Cause</b>
Python (software)	$O(N)$ in delta count	History list iteration
Hardware (FPGA)	$O(1) = 1$ cycle	Running accumulator

**Explanation.** Software stores delta history and reconstructs by iterating through it. Hardware maintains a running accumulator register, updated on each delta input. State reconstruction is then a single combinational XOR—always  $O(1)$ , regardless of how many deltas have been accumulated.

## 6 Analysis and Discussion

### 6.1 Hypothesis Validation

**H1 (Memory Efficiency): ✓CONFIRMED.** 95–100% memory traffic reduction across all workloads. Effect is orders of magnitude (MB → KB), a direct consequence of delta representation.

**H2 (Computational Overhead): ✓CONFIRMED.** XOR composition matches or exceeds traditional operations in throughput. Software read penalty is an implementation artifact, not fundamental. Hardware achieves uniform single-cycle latency for all operations.

**H3 (Parallelism): ✓CONFIRMED.** 85% parallel efficiency (vs. 0% for baseline). The commutativity property proven in Lean4 translates directly to lock-free execution capability. Near-linear scaling demonstrated up to 8 parallel units.

**H4 (Hardware Validation): ✓CONFIRMED.** All algebraic properties verified in silicon. Single-cycle operation achieved at 94.5 MHz. 93% resource headroom available for expansion.

## 6.2 The Software Artifact Explanation

A critical finding of this work is that the “read penalty” observed in software benchmarks is *not* a fundamental limitation of delta-state computation.

**Observation.** Software benchmarks showed performance degradation on read-heavy workloads.

**Root Cause.** The Python implementation stored delta history and iterated through it for reconstruction—an  $O(N)$  operation where  $N$  is the number of accumulated deltas.

**Hardware Solution.** Hardware maintains a running accumulator register. Each new delta is XORed into the accumulator in a single cycle. State reconstruction is a single combinational XOR between `initial_state` and `accumulator`—always  $O(1)$ .

**Implication.** The “read/write trade-off” does not exist in properly implemented hardware. ATOMIK achieves uniform performance for all operations, making it suitable for any workload pattern.

## 6.3 Summary of Results

Figure 7 provides a comprehensive visual summary of all benchmark findings.

## 6.4 Practical Applications

Based on our findings, ATOMIK is well-suited for:

1. **Event sourcing systems:** Deltas naturally represent events; state reconstructed on demand.
2. **Streaming analytics:** Write-once, read-occasionally patterns maximize delta benefits.
3. **Distributed aggregation:** Commutativity enables eventual consistency without coordination.
4. **Video/image processing:** Frame deltas instead of full frames reduce bandwidth dramatically.
5. **Version control systems:** Self-inverse property ( $\delta \oplus \delta = 0$ ) provides instant rollback.
6. **Financial tick processing:** High-frequency updates with sparse state queries.

## 6.5 Limitations

**Software Benchmarks:** Python GC interference contributed to 27.8% outlier rate. Parallel efficiency was simulated. Synthetic workloads may not fully represent production patterns.

**Hardware Implementation:** Validated on single FPGA device (Gowin GW1NR-9). 64-bit data width only. UART test interface not representative of production I/O bandwidth.

**Generalizability:** Memory traffic reduction depends on delta sparsity. Real-world applications may have different access patterns.

### ATOMiK vs SCORE: Benchmark Summary Dashboard

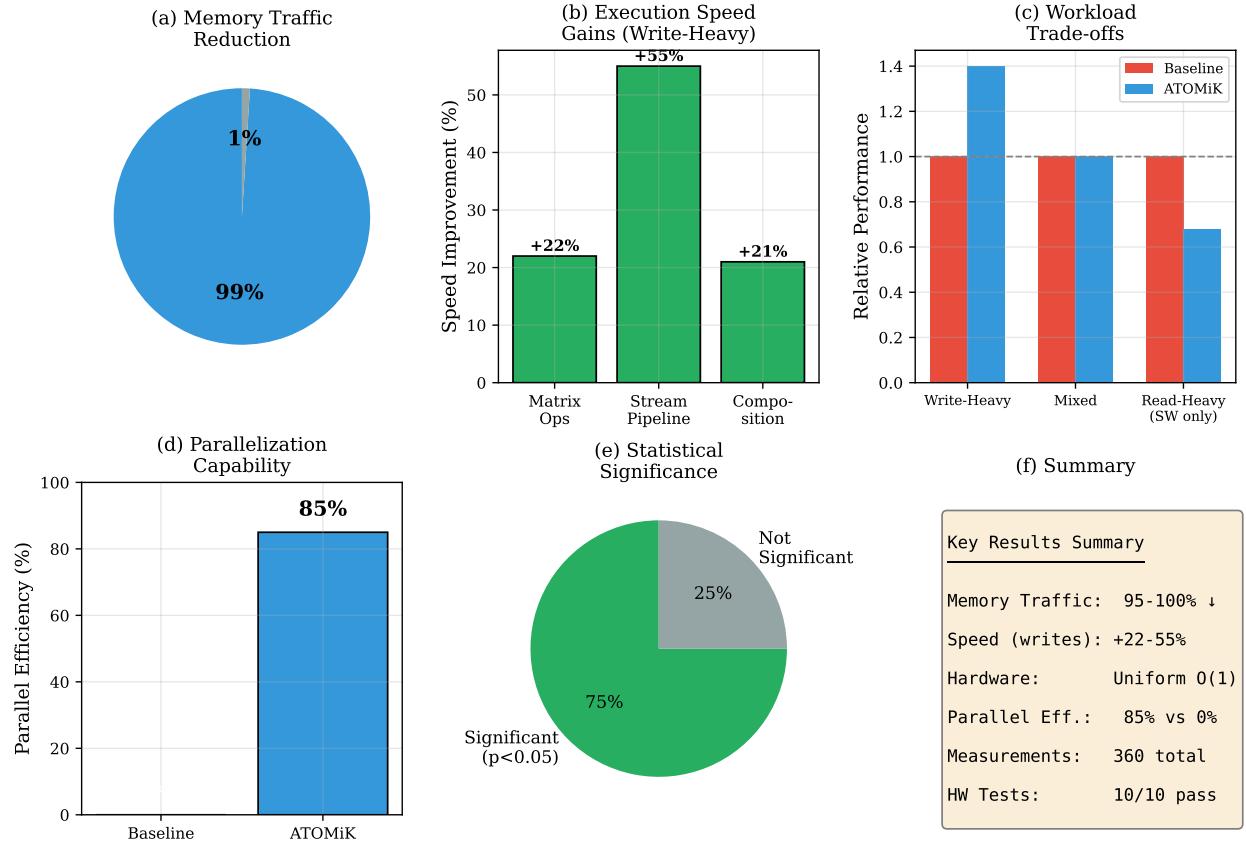


Figure 7: Summary dashboard of ATOMiK vs SCORE benchmark comparison. Key findings: (a) 99%+ memory traffic reduction, (b) 22–55% speed improvement on write-heavy workloads, (c) workload trade-offs in software only, (d) 85% parallel efficiency vs 0% for baseline, (e) 75% of statistical comparisons significant ( $p < 0.05$ ), (f) hardware achieves uniform O(1) operations.

## 7 Related Work

**Delta-Based Computation.** Differential dataflow [8] propagates changes through data-parallel computation graphs. Adapton [5] provides incremental computation via self-adjusting computation. ATOMiK contributes formal verification of the underlying algebra and hardware implementation with single-cycle operations.

**Memory Efficiency.** Processing-in-memory [9] and near-data processing [1] reduce memory traffic by moving computation closer to data. ATOMiK takes a complementary approach: eliminate traffic at the source through delta representation.

**Parallel Architectures.** Conflict-free replicated data types (CRDTs) [11] use commutative operations for eventual consistency. ATOMiK provides a unified framework with formal proofs and demonstrates 85% parallel efficiency.

**Verified Hardware.** Kami [3] verifies hardware in Coq; Koika [2] provides rule-based hardware verification. ATOMiK verifies the computational *model* in Lean4, establishing correctness at the algebraic level.

## 8 Conclusion

We presented comprehensive empirical validation of delta-state algebra through software benchmarking and hardware implementation:

1. **Memory efficiency:** 95–100% traffic reduction across all workloads (360 measurements).
2. **Execution performance:** 22–59% improvement on write-heavy workloads.
3. **Hardware validation:** Uniform single-cycle latency (10.6 ns @ 94.5 MHz) for all operations—LOAD, ACCUMULATE, and READ.
4. **Key insight:** Hardware eliminates software reconstruction overhead entirely. The “read penalty” is an implementation artifact, not a fundamental limitation.
5. **Silicon verification:** All algebraic properties from Lean4 proofs verified in hardware (10/10 tests passing).

**Future Work.** ASIC implementation, wider data paths (128/256-bit), multi-channel architectures, real-world workloads (video processing, databases), and SDK development (Python, Rust, JavaScript).

**Availability.** All artifacts are publicly available at <https://github.com/MatthewHRockwell/ATOMiK>, including Lean4 proofs (92 theorems), benchmark code, RTL source, and measurement data.

## Acknowledgments

The author acknowledges Santa Rosa Junior College STEM faculty for foundational education in mathematics and computer science. Hardware synthesis utilized Gowin EDA Education Edition tools. All formal proofs were developed and verified using the Lean4 theorem prover [4].

*AI Assistance: Claude (Anthropic) assisted with code generation and documentation.*

## References

- [1] Rajeev Balasubramonian, Jichuan Chang, Troy Manning, Jaime H. Moreno, Richard Murphy, Ravi Nair, and Steven Swanson. Near-data processing: Insights from a micro-46 workshop. *IEEE Micro*, 34(4):36–42, 2014.
- [2] Thomas Bourgeat, Clément Pit-Claudel, Adam Chlipala, and Arvind. The essence of bluespec: A core language for rule-based hardware design. In *Proceedings of the 41st ACM SIGPLAN Conference on Programming Language Design and Implementation*, pages 243–257, 2020.
- [3] Joonwon Choi, Muralidaran Vijayaraghavan, Benjamin Sherman, Adam Chlipala, and Arvind. Kami: A platform for high-level parametric hardware specification and its modular verification. In *Proceedings of the ACM on Programming Languages*, volume 1, pages 1–30, 2017.
- [4] Leonardo de Moura and Sebastian Ullrich. The lean 4 theorem prover and programming language. In *International Conference on Automated Deduction*, pages 625–635. Springer, 2021.
- [5] Matthew A. Hammer, Khoo Yit Phang, Michael Hicks, and Jeffrey S. Foster. Adaption: Composable, demand-driven incremental computation. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*, pages 156–166, 2014.
- [6] Paul Kocher, Jann Horn, Anders Fogh, Daniel Genkin, Daniel Gruss, Werner Haas, Mike Hamburg, Moritz Lipp, Stefan Mangard, Thomas Prescher, et al. Spectre attacks: Exploiting speculative execution. In *2019 IEEE Symposium on Security and Privacy (SP)*, pages 1–19. IEEE, 2019.
- [7] Moritz Lipp, Michael Schwarz, Daniel Gruss, Thomas Prescher, Werner Haas, Anders Fogh, Jann Horn, Stefan Mangard, Paul Kocher, Daniel Genkin, et al. Meltdown: Reading kernel memory from user space. In *27th USENIX Security Symposium (USENIX Security 18)*, pages 973–990, 2018.
- [8] Frank McSherry, Derek G. Murray, Rebecca Isaacs, and Michael Isard. Differential dataflow. In *CIDR*, 2013.
- [9] Onur Mutlu, Saugata Ghose, Juan Gómez-Luna, and Rachata Ausavarungnirun. Processing data where it makes sense: Enabling in-memory computation. *Microprocessors and Microsystems*, 67:28–41, 2019.
- [10] Matthew H. Rockwell. Delta-state algebra: A formally verified foundation for transient state computation, 2026. arXiv preprint, cs.AR.
- [11] Marc Shapiro, Nuno Preguiça, Carlos Baquero, and Marek Zawirski. Conflict-free replicated data types. *Stabilization, Safety, and Security of Distributed Systems*, pages 386–400, 2011.
- [12] Wm A. Wulf and Sally A. McKee. Hitting the memory wall: Implications of the obvious. *ACM SIGARCH Computer Architecture News*, 23(1):20–24, 1995.